

Benchmarking ImageNet Transfer Learning for Restricted Dog Image Classification

Ronan Downes

MSc in Data Analytics, CCT College Dublin


Supervisor: Sam Weiss

Research Aim

To benchmark six ImageNet-pretrained CNNs for restricted-versus-unrestricted dog-image classification and evaluate the effect of limited fine-tuning.

Dataset and Experimental Design

- Stanford Dogs image dataset: **30 breeds**.
- 6 restricted** and **24 unrestricted** breeds.
- Fixed training, validation and held-out test splits.
- Six ImageNet-pretrained CNNs benchmarked as frozen feature extractors.
- The selected model was fine-tuned by unfreezing its final 20 layers.
- Evaluation used accuracy, precision, recall, F1, ROC-AUC, PR-AUC, MCC and confusion counts.



Six Benchmark Architectures

VGG16 · ResNet50 · InceptionV3
Xception · InceptionResNetV2 · NASNetMobile

Reproducibility


- Data preparation
- Frozen-model benchmark
- Selected-model fine-tuning
- Grad-CAM explanation

The archive retains the four notebooks, figures, result tables, predictions, checkpoints and complete $\text{\LaTeX}$ source.

Ethical Boundary

High predictive performance does not justify autonomous enforcement.

Frozen Model Benchmark



Best frozen model: InceptionResNetV2

97.08% accuracy · 98.42% recall · 92.57% F1


Fine-Tuning Result

Metric	Frozen	Fine-tuned
Accuracy	97.08%	98.25%
Precision	87.38%	93.88%
Recall	98.42%	96.84%
F1	92.57%	95.34%
MCC	0.9101	0.9428


Key trade-off: 15 fewer false positives, but 3 additional false negatives.

Fine-tuning improved overall balance and precision, but slightly reduced recall.

Fine-Tuned Confusion Matrix




Model Explainability



True positive: German Shepherd

Attention concentrated around the head and ears, which contributed strongly to the restricted prediction.



False positive: Weimaraner

Body shape, pose, coat contrast and contextual structure may have contributed to the incorrect restricted prediction.

Interpretive boundary: Grad-CAM heatmaps are post-hoc visual diagnostics. They indicate regions associated with a prediction, but they do not provide causal, behavioural or legal explanations.

Conclusions

- Six frozen ImageNet CNNs were benchmarked under one pipeline.
- InceptionResNetV2 produced the strongest frozen result.
- Fine-tuning improved precision and overall balance but reduced recall.
- The classifier cannot establish legal identity, ancestry, behaviour or dangerousness.