

Restricted Dog Breed Classification Using CNN Transfer Learning

Ronan Downes

MSc in Data Analytics

CCT College Dublin

Supervisor: Sam Weiss

April 2025

Research problem

- Benchmark six pretrained CNNs
- Binary dog-image classification
- Irish restricted-breed context
- Stanford Dogs dataset
- Controlled model comparison

Research problem

- Benchmark six pretrained CNNs
- Binary dog-image classification
- Irish restricted-breed context
- Stanford Dogs dataset
- Controlled model comparison

Research problem

- Benchmark six pretrained CNNs
- Binary dog-image classification
- Irish restricted-breed context
- Stanford Dogs dataset
- Controlled model comparison

Research problem

- Benchmark six pretrained CNNs
- Binary dog-image classification
- Irish restricted-breed context
- Stanford Dogs dataset
- Controlled model comparison

Research problem

- Benchmark six pretrained CNNs
- Binary dog-image classification
- Irish restricted-breed context
- Stanford Dogs dataset
- Controlled model comparison

Research questions

1. How do six frozen CNNs compare?
2. Does fine-tuning improve performance?
3. What does Grad-CAM reveal?
4. What limits real-world use?

Research questions

1. How do six frozen CNNs compare?
2. Does fine-tuning improve performance?
3. What does Grad-CAM reveal?
4. What limits real-world use?

Research questions

1. How do six frozen CNNs compare?
2. Does fine-tuning improve performance?
3. What does Grad-CAM reveal?
4. What limits real-world use?

Research questions

1. How do six frozen CNNs compare?
2. Does fine-tuning improve performance?
3. What does Grad-CAM reveal?
4. What limits real-world use?

Stanford Dogs Restricted-breeds

Represented

Rottweiler

German Shepherd

Doberman Pinscher

Bull Mastiff

Rhodesian Ridgeback

Staffordshire Bull Terrier

Not represented

American Pit Bull Terrier

English Bull Terrier

Japanese Akita

Japanese Tosa

Bandog

Dataset construction

6

restricted breeds

vs.

24

unrestricted breeds

24 breeds reproducibly selected
from the remaining 114 Stanford categories

- Broader negative class
- Manageable computational scope
- Fixed data splits

Dataset construction

6

restricted breeds

vs.

24

unrestricted breeds

24 breeds reproducibly selected
from the remaining 114 Stanford categories

- Broader negative class
- Manageable computational scope
- Fixed data splits

Dataset construction

6

restricted breeds

vs.

24

unrestricted breeds

24 breeds reproducibly selected
from the remaining 114 Stanford categories

- Broader negative class
- Manageable computational scope
- Fixed data splits


Experimental pipeline




Experimental pipeline




Experimental pipeline




Experimental pipeline




Experimental pipeline



Experimental pipeline



Experimental pipeline



ILSVRC Winners Record

Year	Architecture	Key idea
2010	SIFT/LBP + SVM	Hand-crafted features
2011	Fisher vectors	Best pre-CNN
2012	AlexNet	Deep-CNN breakthrough
2014	VGG16	Uniform deep CNN
2015	ResNet50	Residual connections
2015	InceptionV3	Multi-scale processing
2016	Xception	Separable convolutions
2016	InceptionResNetV2	Inception plus residual
2017	NASNetMobile	Neural architecture search

Six benchmarked CNN models

- **VGG16**

Uniform deep baseline

- **ResNet50**

Residual connections

- **InceptionV3**

Multi-scale processing

- **Xception**

Separable convolutions

- **InceptionResNetV2**

Multi-scale residual model

- **NASNetMobile**

Architecture-search model

Six benchmarked CNN models

- **VGG16**

Uniform deep baseline

- **ResNet50**

Residual connections

- **InceptionV3**

Multi-scale processing

- **Xception**

Separable convolutions

- **InceptionResNetV2**

Multi-scale residual model

- **NASNetMobile**

Architecture-search model

Six benchmarked CNN models

- **VGG16**

Uniform deep baseline

- **ResNet50**

Residual connections

- **InceptionV3**

Multi-scale processing

- **Xception**

Separable convolutions

- **InceptionResNetV2**

Multi-scale residual model

- **NASNetMobile**

Architecture-search model

Six benchmarked CNN models

- **VGG16**

Uniform deep baseline

- **ResNet50**

Residual connections

- **InceptionV3**

Multi-scale processing

- **Xception**

Separable convolutions

- InceptionResNetV2

Multi-scale residual model

- NASNetMobile

Architecture-search model

Six benchmarked CNN models

- **VGG16**

Uniform deep baseline

- **ResNet50**

Residual connections

- **InceptionV3**

Multi-scale processing

- **Xception**

Separable convolutions

- **InceptionResNetV2**

Multi-scale residual model

- **NASNetMobile**

Architecture-search model

Six benchmarked CNN models

- **VGG16**

Uniform deep baseline

- **ResNet50**

Residual connections

- **InceptionV3**

Multi-scale processing

- **Xception**

Separable convolutions

- **InceptionResNetV2**

Multi-scale residual model

- **NASNetMobile**

Architecture-search model

Frozen Benchmark Approach

- ImageNet weights
 - Frozen feature extractors
 - Common classification head
 - Class weighting
 - Early stopping
 - Best validation checkpoint
- Multiple metrics—not accuracy alone

Frozen Benchmark Approach

- ImageNet weights
 - Frozen feature extractors
 - Common classification head
 - Class weighting
 - Early stopping
 - Best validation checkpoint
- Multiple metrics—not accuracy alone

Frozen Benchmark Approach

- ImageNet weights
- Frozen feature extractors
- Common classification head
- Class weighting
- Early stopping
- Best validation checkpoint

Multiple metrics—not accuracy alone

Frozen Benchmark Approach

- ImageNet weights
- Frozen feature extractors
- Common classification head
- Class weighting
- Early stopping
- Best validation checkpoint

Multiple metrics—not accuracy alone

Frozen Benchmark Approach

- ImageNet weights
- Frozen feature extractors
- Common classification head
- Class weighting
- Early stopping
- Best validation checkpoint

Multiple metrics—not accuracy alone

Frozen Benchmark Approach

- ImageNet weights
- Frozen feature extractors
- Common classification head
- Class weighting
- Early stopping
- Best validation checkpoint

Multiple metrics—not accuracy alone

Frozen Benchmark Approach

- ImageNet weights
- Frozen feature extractors
- Common classification head
- Class weighting
- Early stopping
- Best validation checkpoint


Multiple metrics—not accuracy alone

Frozen Benchmark Results Table

CNN	Acc.	Prec.	Recall	F1
VGG16	87.4	60.3	92.1	72.9
ResNet50	91.7	70.9	93.7	80.7
InceptionV3	96.1	84.4	96.8	90.2
Xception	96.0	83.7	97.4	90.0
InceptionResNetV2	96.8	86.9	97.4	91.8
NASNetMobile	94.8	79.9	96.3	87.4

Percentages shown; bold and shading are selected model.

Frozen Benchmark Results Chart



Fine-tuning strategy

Task adaptation without destabilising pretrained features

- Last 20 layers unfrozen
- Learning rate: 10^{-5}
- Maximum 15 epochs
- Class weighting retained
- Early stopping
- Best validation checkpoint

Fine-tuning strategy

Task adaptation without destabilising pretrained features

- Last 20 layers unfrozen
- Learning rate: 10^{-5}
- Maximum 15 epochs
- Class weighting retained
- Early stopping
- Best validation checkpoint

Fine-tuning strategy

Task adaptation without destabilising pretrained features

- Last 20 layers unfrozen
- Learning rate: 10^{-5}
- Maximum 15 epochs
- Class weighting retained
- Early stopping
- Best validation checkpoint

Fine-tuning strategy

Task adaptation without destabilising pretrained features

- Last 20 layers unfrozen
- Learning rate: 10^{-5}
- Maximum 15 epochs
- Class weighting retained
- Early stopping
- Best validation checkpoint

Fine-tuning strategy

Task adaptation without destabilising pretrained features

- Last 20 layers unfrozen
- Learning rate: 10^{-5}
- Maximum 15 epochs
- Class weighting retained
- Early stopping
- Best validation checkpoint

Fine-tuning strategy


Task adaptation without destabilising pretrained features

- Last 20 layers unfrozen
- Learning rate: 10^{-5}
- Maximum 15 epochs
- Class weighting retained
- Early stopping
- Best validation checkpoint


Fine-tuning results

Metric	Frozen	Fine-tuned
Accuracy	96.79%	98.25%
Precision	86.85%	93.88%
Recall	97.37%	96.84%
F1	91.81%	95.34%
MCC	0.9005	0.9428

Fewer false positives
Slightly more false negatives





Learning behaviour




Best validation checkpoint retained


Grad-CAM explainability



True positive



True negative



False positive

False negative

Diagnostic evidence—not proof of causal reasoning.

Project Contribution

- Reproducible pipeline
- Six-CNN benchmarks
- Multi-metric model selection
- Fine-tuning analysis
- Error trade-off analysis
- Grad-CAM explainability

Project Contribution

- Reproducible pipeline
- Six-CNN benchmarks
- Multi-metric model selection
- Fine-tuning analysis
- Error trade-off analysis
- Grad-CAM explainability

Project Contribution

- Reproducible pipeline
- Six-CNN benchmarks
- Multi-metric model selection
- Fine-tuning analysis
- Error trade-off analysis
- Grad-CAM explainability

Project Contribution

- Reproducible pipeline
- Six-CNN benchmarks
- Multi-metric model selection
- Fine-tuning analysis
- Error trade-off analysis
- Grad-CAM explainability

Project Contribution

- Reproducible pipeline
- Six-CNN benchmarks
- Multi-metric model selection
- Fine-tuning analysis
- Error trade-off analysis
- Grad-CAM explainability

Project Contribution

- Reproducible pipeline
- Six-CNN benchmarks
- Multi-metric model selection
- Fine-tuning analysis
- Error trade-off analysis
- Grad-CAM explainability

IPYNB Notebook Execution

- | | |
|------------------------------|--------------|
| 1. 01_Data_Preparation.ipynb | CPU |
| 2. 02_Benchmarking.ipynb | GPU required |
| 3. 03_Fine_Tuning.ipynb | GPU required |
| 4. 04_Gradcam.ipynb | CPU |

Local hardware did not provide a suitable GPU, so Google Colab was used to accelerate benchmarking and fine-tuning with a **GPU hardware accelerator** (NVIDIA T4 when available). Free-tier compute limits occasionally restricted GPU access, so some runs were completed during off-peak periods.

IPYNB Notebook Execution

- | | |
|------------------------------|--------------|
| 1. 01_Data_Preparation.ipynb | CPU |
| 2. 02_Benchmarking.ipynb | GPU required |
| 3. 03_Fine_Tuning.ipynb | GPU required |
| 4. 04_Gradcam.ipynb | CPU |

Local hardware did not provide a suitable GPU, so Google Colab was used to accelerate benchmarking and fine-tuning with a **GPU hardware accelerator** (NVIDIA T4 when available). Free-tier compute limits occasionally restricted GPU access, so some runs were completed during off-peak periods.

IPYNB Notebook Execution

- | | |
|------------------------------|--------------|
| 1. 01_Data_Preparation.ipynb | CPU |
| 2. 02_Benchmarking.ipynb | GPU required |
| 3. 03_Fine_Tuning.ipynb | GPU required |
| 4. 04_Gradcam.ipynb | CPU |

Local hardware did not provide a suitable GPU, so Google Colab was used to accelerate benchmarking and fine-tuning with a **GPU hardware accelerator** (NVIDIA T4 when available). Free-tier compute limits occasionally restricted GPU access, so some runs were completed during off-peak periods.

IPYNB Notebook Execution

- | | |
|------------------------------|--------------|
| 1. 01_Data_Preparation.ipynb | CPU |
| 2. 02_Benchmarking.ipynb | GPU required |
| 3. 03_Fine_Tuning.ipynb | GPU required |
| 4. 04_Gradcam.ipynb | CPU |

Local hardware did not provide a suitable GPU, so Google Colab was used to accelerate benchmarking and fine-tuning with a **GPU hardware accelerator** (NVIDIA T4 when available). Free-tier compute limits occasionally restricted GPU access, so some runs were completed during off-peak periods.

Limitations?

Deployment?

Further Questions?

Conclusions

1. Transfer learning performed strongly
2. InceptionResNetV2 ranked first
3. Fine-tuning improved overall balance
4. Fine-tuning changed the error trade-off
5. Current limitations do not justify real-world deployment

Conclusions

1. Transfer learning performed strongly
2. InceptionResNetV2 ranked first
3. Fine-tuning improved overall balance
4. Fine-tuning changed the error trade-off
5. Current limitations do not justify real-world deployment

Conclusions

1. Transfer learning performed strongly
2. InceptionResNetV2 ranked first
3. Fine-tuning improved overall balance
4. Fine-tuning changed the error trade-off
5. Current limitations do not justify real-world deployment

Conclusions

1. Transfer learning performed strongly
2. InceptionResNetV2 ranked first
3. Fine-tuning improved overall balance
4. Fine-tuning changed the error trade-off
5. Current limitations do not justify real-world deployment

Conclusions

1. Transfer learning performed strongly
2. InceptionResNetV2 ranked first
3. Fine-tuning improved overall balance
4. Fine-tuning changed the error trade-off
5. Current limitations do not justify real-world deployment